

# Introduction to Computer Organization

DIS 1H – Week 3

Slides modified from UT WANG

# Logistics

Extra office hours – **Week 4**

Monday 2:00 P.M. – 4:00 P.M.

and

Friday 12:00 P.M. – 2:00 P.M.

**Week 5** – No discussion on May 5, 2017

# Agenda

- Switch statement
- Arrays and Structures
- Procedures
- Midterm review
- Arrays example
- More Struct and Union
- gdb

### Practice Problem 3.37

Consider the following source code, where  $M$  and  $N$  are constants declared with `#define`:

```
1  int mat1[M][N];
2  int mat2[N][M];
3
4  int sum_element(int i, int j) {
5      return mat1[i][j] + mat2[j][i];
6  }
```

In compiling this program, GCC generates the following assembly code:

```
    i at %ebp+8, j at %ebp+12
1    movl    8(%ebp), %ecx
2    movl    12(%ebp), %edx
3    leal    0(,%ecx,8), %eax
4    subl    %ecx, %eax
5    addl    %edx, %eax
6    leal    (%edx,%edx,4), %edx
7    addl    %ecx, %edx
8    movl    mat1(,%eax,4), %eax
9    addl    mat2(,%edx,4), %eax
```

Use your reverse engineering skills to determine the values of  $M$  and  $N$  based on this assembly code.

---

### Solution to Problem 3.37 (page 236)

This problem requires you to work through the scaling operations to determine the address computations, and to apply Equation 3.1 for row-major indexing. The first step is to annotate the assembly code to determine how the address references are computed:

1	movl	8(%ebp), %ecx	<i>Get i</i>
2	movl	12(%ebp), %edx	<i>Get j</i>
3	leal	0(,%ecx,8), %eax	$8*i$
4	subl	%ecx, %eax	$8*i - i = 7*i$
5	addl	%edx, %eax	$7*i + j$
6	leal	(%edx,%edx,4), %edx	$5*j$
7	addl	%ecx, %edx	$5*j + i$
8	movl	mat1(,%eax,4), %eax	$mat1[7*i + j]$
9	addl	mat2(,%edx,4), %eax	$mat2[5*j + i]$

We can see that the reference to matrix mat1 is at byte offset  $4(7i + j)$ , while the reference to matrix mat2 is at byte offset  $4(5j + i)$ . From this, we can determine that mat1 has 7 columns, while mat2 has 5, giving  $M = 5$  and  $N = 7$ .

# Structs

```
struct s {  
    char c1;  
    int i;  
    char c2;  
    int j;  
};
```

- What's the problem with this struct?

# Structs

```
struct s {  
    char c1;  
    int i;  
    char c2;  
    int j;  
};
```

- Say an instance of the struct begins at 0x10. Then c1 is at address 0x10. However, 'i' cannot be at address 0x11 (it needs to be 4-aligned). As a result, we need 3 bytes of padding.

# Structs

- This is a waste of space! There will be 3 bytes of padding after c1 and 3 bytes of padding after c2, meaning that this struct will take up 16 bytes when really it only needs 10.



# Structs

- Two common struct ordering guidelines (which could be at odds):
  1. Place the most commonly used data type first.
  2. Place the elements in descending order of size (ie largest first)
- Why?

# Structs

- 1.
- Memory references are expensive (ex. (%eax))... but memory references with an offset are more expensive (ex. 8(%eax))
- Chances are, you'll be referring to the struct by a pointer to the beginning of the struct, which means that dereferencing the pointer without an offset will point to the first element.

# Structs

- 2.
- If the elements with larger sizes are first, that means there will be less of a need for padding.
- For example, consider struct s, except with the first two elements swapped:

```
struct s {  
    int i;  
    char c1;  
    char c2;  
    int j;  
};
```

# Structs

- 2.
- ```
struct s {  
    int i;  
    char c1;  
    char c2;  
    int j;  
};
```
- Now, we need 2 bytes of padding between c2 and j for a total of 12 bytes.

# Structs

- Because each internal element must follow their own alignment rules, the alignment of the struct must be equal to the strictest of the elements within a struct.
- But wait...

# Structs

- Consider:

```
struct s {  
    char c;  
    int i;  
};
```

- Because int i is aligned by 4, instances of struct s must be aligned by 4.
- There must also be 3 bytes of padding between c and i, meaning a total size of 8.

# Structs

- Thus, a possible placement of (struct s s1) where s1.c = 0xFF and s1.i = 0x33221100 is the following:

|          |      |      |      |      |      |      |      |      |
|----------|------|------|------|------|------|------|------|------|
| Address: | 0x10 | 0x11 | 0x12 | 0x13 | 0x14 | 0x15 | 0x16 | 0x17 |
| Value:   | 0xFF | 0XX  | 0XX  | 0XX  | 0x00 | 0x11 | 0x22 | 0x33 |

- Where s begins at 0x10.
- This is how we meet the alignment requirements of each individual item

# Unions

- Like structs except all of the values begin at the same address.
- union s {
  - short s;
  - char c;
  - };
- This means that in a union that contains several values, only one of them is likely to be meaningful and assigning one term a value will trample other terms.



# Unions

- union s {
  - short s;
  - char c;
  - };
- union s foo;
- Say foo begins at 0x10.
- foo.s will be located in addresses 0x10 and 0x11
- foo.c will be located in address 0x10.

# Unions

- union s {
- short s;
- char c;
- };
- union s foo;
- foo.s = 0xFFFF;
- foo.c = 0;
- printf(“%hx\n”, foo.s) => FF00

# **`gdb` - Debugger**

`(gdb) break <function_name>`

`(gdb) run (gdb)`

`(gdb) stepi`

`(gdb) stepn`

`(gdb) info registers`

`(gdb) disassemble`

All variants of “print”

Cheat Sheet – `gdbnotes.pdf`

Example – `gdb.pdf`